

completeness  
o Space complexity  
o Correctness  
o An algorithm is complete if it will find a solution when one exists and correctly report failure (no solution) when it does not  
o An algorithm is optimal if it finds a solution with the lowest path cost among all solutions (i.e., path cost optimal)  
Actions:  $s_i \rightarrow A$   
Function that returns the possible actions,  $A(s_1, \dots, a_k)$ , at a given state  $s_i$   
Initial state  
Transition model,  $T(s, a, s')$   
Function that returns the intermediate state transitioned to,  $s'$ , when action  $a_j$  is applied at state  $s_i$   
Goal test,  $isGoal(s) \rightarrow \{0, 1\}$   
Function that returns 1 if given state  $s_i$  is a goal state, else returns 0  
Action costs, cost:  $(s_i, a_j, s')$   $\rightarrow v$   
Function that returns cost,  $v$ , of taking the action  $a_j$  at state  $s_i$  to reach the intermediate state  $s'$   
Generally assume costs  $>= 0$

general search algo

```
Function TreeSearch(initial_state, actions, T, isGoal, cost):  
    frontier = (Node(initial_state, NULL))  
    while frontier not empty:  
        current = frontier.pop()  
        if isGoal(current.state): return current.getPath()  
        for a in actions(current.state):  
            successor = Node(T(current.state, a), current)  
            frontier.push(successor)  
    return failure
```

Default assumptions on search spaces:

- finite
- finite contains a solution
- infinite
- All actions costs  $\geq 0$

Exercise questions:

- Show that when  $b$  is finite, BFS is incomplete. Show that when  $b$  is infinite, but  $a$  is infinite AND there is no solution, BFS is incomplete. Under what (other) cases is BFS optimal? Why is the complexity of BFS  $O(b^d)$  in this scenario?

Exercise questions:

- Show that DFS under BFS completeness assumptions is incomplete.
- Can space complexity be improved?

Early goal test for BFS:

Updating goal test on pushing to frontier instead of popping from frontier will prevent this with no change in the complexity or the BFS solution. However, it should be noted that a different valid path may be output

Updated default assumptions on search spaces:

- All action costs  $\geq 0$

Exercise questions:

- Can you show that UCS under BFS completeness assumptions is incomplete when action costs are not  $\geq 0$  (e.g., when  $c > 0$ )?
- Why is the complexity of UCS  $O(b^d)$ ?
- Why is UCS optimal?
- Would UCS still be optimal with an early goal test?

UCS:

Frontier: Priority Queue<sup>1</sup>

Time Complexity:  $O(b^d)$

Space Complexity:  $O(b^d)$

Complete: Yes<sup>2</sup>

Optimal: Yes

$b$ : branching factor  
 $c$ :  $1 + \lceil c^* / \epsilon \rceil$ , where  $c^*$  is the optimal path cost and  $\epsilon$  is some small positive constant  
1: prioritizing lower path cost,  $g(s)$ , where  $g(s)$  = path cost of the path taken to reach  $s$   
2: requires BFS completeness criteria and that action costs  $> \epsilon > 0$   
Note: determining path cost for each node is  $O(1)$  since we store the current path cost in each node

DFS:

Frontier: Stack

Time Complexity:  $O(b^d)$

Space Complexity:  $O(bm)$

Complete: No<sup>1</sup>

Optimal: No

$b$ : branching factor  
 $c$ : maximum depth  
1: DFS may be incomplete even if a solution exists

Depth-Limit Search (DLS)

DFS with a depth limit,  $\ell$

- Search only up to depth  $\ell$
- Assume no actions may be taken from nodes at depth  $\ell$

Same guarantees as DFS with  $\ell$  in place of  $\infty$

- Time complexity:  $O(b^\ell)$
- Space complexity:  $O(b \cdot \ell)$
- Complete: No
- Optimal: No

Iterative Deepening Search (IDS)

Example,  $b = 10$  and  $d = 5$

- $N_{total} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead = 11%

IDS properties

- Time:  $O(b^d)$
- Space:  $O(bd)$
- Complete: Yes (if  $b$  is finite and  $d$  is finite or contains solution) – same as BFS
- Optimal: No (optimal if costs uniform (and some other cases)) – same as BFS

Summary tree search

Criterion	BFS <sup>1</sup>	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No <sup>3</sup>	No	Yes <sup>1</sup>
Optimal Cost?	No <sup>4</sup>	Yes	No	No	No <sup>4</sup>
Time	$O(b^d)$	$O(b^{1+\lceil c^*/\epsilon \rceil})$	$O(b^d)$	$O(b^d)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+\lceil c^*/\epsilon \rceil})$	$O(bm)$	$O(b \cdot d)$	$O(b \cdot d)$

- Complete if  $b$  is finite and either has a solution or  $a$  is finite
  - Complete if all actions costs are  $> \epsilon > 0$
  - DFS is incomplete unless the search space is finite – i.e., when  $b$  is finite and  $a$  is finite
  - Cost optimal if action costs are all identical (and several other cases)
- Recall that an Early Goal Test is assumed for BFS
  - UCS must perform a Late Goal Test to be optimal (this also accounts for the  $+1$  in the index of its complexity)
  - DFS is incomplete (even under 1 – note the “or”); if a solution exists, it may infinitely traverse a path without a solution
  - DFS space complexity may be improved to  $O(m)$  with backtracking (similar for DLS and IDS)

Tree search allows revisits, cyclic graph, incomplete  
Summary graph search

Criterion	BFS <sup>1</sup>	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No <sup>3</sup>	No	Yes <sup>1</sup>
Optimal Cost?	No <sup>4</sup>	Yes	No	No	No <sup>4</sup>
Time					
Space					

- Complete if  $b$  is finite and either has a solution or  $a$  is finite
  - Complete if all actions costs are  $> \epsilon > 0$
  - DFS is incomplete unless the search space is finite – i.e., when  $b$  is finite and  $a$  is finite
  - Cost optimal if action costs are all identical (and several other cases)
- Time and space complexities are now bounded by the size of the search space – i.e., the number of vertices (nodes) and edges,  $(V) + (E)$
- Note that we do not need to check for cheaper paths under graph search for BFS and DFS since costs play no part in those algorithms and they cannot guarantee an optimal solution anyway

Admissible paths not ending at a goal are never over estimated  
- non goal,  $n$ ,  $f(n) = g(n) + h(n) \leq g(n) + h^*(n)$   
paths ending at a goal are exact  
- goal  $m$ ,  $f(m) = g(m) + h(m)$ , where  $h(m) = 0$   
if admissible, then  $A^*$  optimal with tree search

Consistency  
 $h(n)$  is consistent, if  $\forall n, n': h(n) \leq \text{cost}(n, n') + h(n')$   
if  $h(n)$  is consistent,  $A^*$  with graph search is optimal  
If,  $h_1(n) \geq h_2(n)$ , then  $h_1$  dominates  $h_2$   
if  $h_1$  is also admissible  
-  $h_1$  must be closer to  $h^*$  than  $h_2$   
-  $h_1$  is usually more efficient than  $h_2$  when used with  $A^*$

Prove that  $A^*$  Search is optimal under graph search v3 with consistent heuristic.  
o Base case: When expanding the start node, it is obvious that  $g(s) = g^*(s)$   
o Inductive case: Assume the inductive hypothesis is true up to  $sk$ .  
So, from  $s_0$  to  $sk$ , the optimal path is found, and each  $g$  at  $sk$  is  $g^*(sk)$   
the optimal path cost from  $s_0$  to  $sk+1$  is  
 $h(sk+1) = g(sk+1) - (A)$   
By consistency, we know that  $sk$  must be popped before  $sk+1$ .  
So, after  $sk$  is popped, we know that when  $sk+1$  is popped, cost to node  $sk+1$  is the minimum of all the ways that it has been reached so far (asserting priority queue based on  $f$ ).  
 $g(sk+1) + h(sk+1) = \text{MIN}(g(sk+1) + h(sk+1), g(sk) + c(sk, sk+1) + h(sk+1))$   
 $g(sk+1) = \text{min}(g(sk+1), g(sk) + c(sk, sk+1))$   
 $\leq g(sk) + \text{cost}(sk, sk+1)$   
 $= g^*(sk) + \text{cost}(sk, sk+1)$   
 $= g^*(sk+1) - (B)$

Graph search V3

```
Function GraphSearchV3(initial_state, actions, T, isGoal, cost):  
    frontier = (Node(initial_state, NULL))  
    reached = {}  
    while frontier not empty:  
        current = frontier.pop()  
        reached.insert(current.state: current)  
        if isGoal(current.state): return current.getPath()  
        for a in actions(current.state):  
            successor = Node(T(current.state, a), current)  
            if successor.state not in reached:  
                frontier.push(successor)  
    return failure
```

Consistent > admissible

Solution: The proof is by induction of  $k(n)$ , which denoted the number of actions required to reach the goal from a node  $n$  to the goal node  $t$ .  
Base case ( $k = 1$ , i.e., the node  $n$  is one step from  $t$ ): Since the heuristic function  $h$  is consistent,  $h(n) \leq c(n, t) + h(t)$ . And since  $h(t) = 0$ ,  $h(n) \leq c(n, t) = h^*(n)$ . Therefore,  $h$  is admissible.  
Induction case: Suppose that our assumption holds for every node that is  $k - 1$  actions away from  $t$ , and let us observe a node  $n$  that is  $k$  actions away from  $t$ ; that is, the least-actions optimal path from  $n$  to  $t$  has  $k > 1$  steps.  
We write the optimal path from  $n$  to  $t$  as:  $n \rightarrow n_1 \rightarrow n_2 \rightarrow \dots \rightarrow nk-1 \rightarrow t$ .  
Since  $h$  is consistent, we have  $h(n) \leq c(n, n_1) + h(n_1)$ .  
Now, note that since  $n_1$  is on a least-cost path from  $t$  to  $n$ , we must have that the path  $n_1 \rightarrow n_2 \rightarrow \dots \rightarrow nk-1 \rightarrow t$  is a minimal-cost path from  $n_1$  to  $t$  as well. By our induction hypothesis, we have  $h(n_1) \leq h^*(n_1)$ .  
Consequently, combining the two inequalities above, we have,  $h(n) \leq c(n, n_1) + h^*(n_1)$ .  
Note that  $h^*(n_1)$  is the cost of the optimal path from  $n_1$  to  $t$ ; by our previous observation (that  $n_1 \rightarrow n_2 \rightarrow \dots \rightarrow nk-1 \rightarrow t$  is an optimal cost path from  $n_1$  to  $t$ ), we have that the cost of the optimal path from  $n$  to  $t$  – i.e.,  $h^*(n)$  – is exactly  $c(n, n_1) + h^*(n_1)$ , which concludes the proof.

admissibility but inconsistent state an admissible heuristic that overestimates the path cost from parent to child (happens when child to goal is understated too much, so  $h(n) > \text{cost}(n \rightarrow n') + h(n')$  ( $h(n')$  understated)

- Completeness assumptions:
- $b$  finite AND ( $a$  finite OR has a solution)
  - All action costs are  $> \epsilon > 0$
- UCS
    - Optimal under
      - Tree search
      - Graph search (version 2 and 3)
    - Popping node  $n \Rightarrow$  optimal path to  $n$  found
  - Greedy Best-First Search
    - Not optimal
    - Incomplete under tree search
- Assume graph search version 1 for Greedy (unless otherwise stated)

Tree vs graph search  
- graph-search algorithms will only explore non-redundant paths  
- only explores nodes that either have not been visited or have been visited but previously via less optimal paths  
- tree search algorithms have no such restrictions; they consider all paths, including redundant ones

Let  $s$  be the initial state,  $n$  be an intermediate state along the optimal path,  $t$  be a suboptimal goal state (i.e., a goal state reached via a suboptimal path), and  $t'$  be the goal along the optimal path.  
An optimal solution implies that  $n$  must be expanded before  $t$ .  
Proof by contradiction:  
- Let us assume that a suboptimal solution is found – i.e., that  $t$  is expanded before  $n$ , which implies that  $f(t) \leq f(n)$ .  
- In other words, given the above frontier, only when  $f(t) < f(n)$  would we expand  $t$  before  $n$ .  
- However, since  $t$  is not on the optimal path but  $t'$  is, we have:  
 $f(t) > g(t) > g(t')$   
 $f(t) > g(t) //$  since  $h(t) = 0$   
 $f(t) > g(n) + p(n, t') //$  where  $p(n, t')$  is the actual cost from  $n$  to  $t'$   
 $f(t) > g(n) + h(n) //$  asserting admissibility  
 $f(t) > f(n) //$  this contradicts (A)  
- Note: we do not consider  $f(t) = f(n)$  since that would mean  $f(t)$  is equally optimal.

- A\* Search**
    - Assuming admissible  $h$ 
      - May not have monotonically increasing path cost
      - Optimal under
        - Tree search
        - Graph search (version 2)
    - Assuming consistent  $h$ 
      - On popping node  $n$ , optimal path to  $n$  found
      - Optimal under
        - Tree search
        - Graph search (version 2 and 3)
- Assume graph search version 2 and late goal testing for UCS and A\* (unless otherwise stated)

$h_1$	$h_2$	$MAX(h_1, h_2)$
Admissible	Admissible	Admissible
Admissible	Inadmissible	Inadmissible
Inadmissible	Inadmissible	Inadmissible

What about consistent heuristics?

- MAX of 2 consistent heuristics is also consistent.
- MIN of 2 consistent heuristics is also consistent.

Try to determine these as an exercise! Might appear in the midterm!

Local Beam Search  
Store  $k$  states instead of 1  
o Hill climbing just stores the current state  
o Beam (window) stores  $k$ 

- General algorithm
- o Begins with  $k$  random starts
- o Each iteration generates successors for each of the  $k$  random start states
- o Repeat with best  $k$  among ALL generated successors unless goal found
- Better than  $k$  parallel random restarts
- o Since best  $k$  among ALL successors taken (not best from each set of successors,  $k$  times)
- Stochastic beam search
- o Original variant may still get stuck in a local cluster
- o Adopt stochastic strategy similar to stochastic hill climbing to increase state diversity

```
Function BeamSearch(random_initial_state, isGoal, actions, T, choose_best_k):  
    frontier = {}  
    for i in 1 to k:  
        s = random_initial_state()  
        if isGoal(s): return s  
        frontier.push(s)  
    while True:  
        new_frontier = {}  
        while frontier not empty:  
            s = frontier.pop()  
            for a in actions(s):  
                new_s = T(s, a)  
                if isGoal(new_s): return new_s  
                new_frontier.push(new_s)  
        frontier = choose_best_k(new_frontier)
```

As the search may loop infinitely, it may be wise to define and use an iteration threshold

No duplicate states on beam

We may also apply stochastic beam search where instead of choosing the best  $k$ , it chooses successors with probability proportional to the successor's value

- Notice that we may allow states with lower values than their parents to be added since choose\_best\_k will simply choose the best  $k$  successor states generated
- We may vary this rule but then may not be able to always ensure we have  $k$  new candidates

Hill-Climbing Algorithm

```
Function HillClimbing(initial_state, eval_fn, highest_valued_successor):  
    current = initial_state  
    while True:  
        neighbour = highest_valued_successor(current)  
        if eval_fn(neighbour) > eval_fn(current): return current  
        current = neighbour
```

Stochastic hill climbing

- Changes the function highest\_valued\_successor(...)
  - Chooses randomly among states with  $f$ -values better than current
  - May take longer to find a solution but sometimes leads to better solutions
- First-choice hill climbing
- Changes the function highest\_valued\_successor(...)
  - Handles high  $b$  by randomly generating successors until one with better  $f$ -value than current is found (instead of generating all possible successors)
  - May even work with infinite  $b$
- Sideways move
- Replaces  $s < w$
  - Allows continuation when  $\text{eval\_fn}(\text{neighbour}) == \text{eval\_fn}(\text{current})$
- Can traverse shoulders / plateaus
- Random-restart hill climbing
- Different algorithm
  - Adds an outer loop which randomly picks a new starting state
  - Keeps attempting random restarts until a solution is found